

TODAY'S HEADLINE

Anthropic sets Claude Code to Auto Mode by default to protect developers fr...

From August 14 Anthropic makes Auto Mode the default permission mode in Claude Code for Pro, Max and Team plans, so a classifier - not a wall of yes/no prompts - decides which commands need review. Anthropic says the classifier caught 89% of dangerous commands in internal tests,

What we cover today.

01 Remote AI Jobs

Location Services Engineer | Maps Platform (Remote in Europe)

02 AI Showcase Opportunities

CUTC: Transform Hackathon

03 YouTube Content Ideas

I Ran Two AI Agents That Talked To Each Other

04 Viral AI on YouTube

OpenAI's Genius Comeback : How Altman Turned a \$27B Loss Into an EMPIRE?

05 Viral Instagram Reels

Untitled

06 Quantum + AI

QuanTiMedAI: Quantum-Enhanced Time-Series Model guided by Agentic AI for Cardiac Arr

07 AI Self-Improvement (RSI)

Inside the Race to Make AI Build Itself - Time Magazine

08 Elon Musk's AI Vision

Grok Imagine 2.0: xAI's image generator is now second only to OpenAI - tbreak.com

09 AI Model Benchmarks

Text / LLM - Best: Claude Opus 4.x / Runner-up: GPT-5.x

10 New AI Tools

Backflip AI turns 3D scans into editable CAD models in minutes instead of hours

11 Indian AI Industry

Government examines Meta's recommendation system for intermediary status

12 Anthropic & Claude Code

Anthropic sets Claude Code to Auto Mode by default to protect developers from bad ap

13 AI Automation & Business

AI is reimaging service companies; Ola Electric narrows loss in Q1

14 Global AI News

AI Data Centers Are Causing Unfathomable Amounts of Air Pollution, and It Gets Worse

15 What People Are Searching For

generative ai meaning

16 Unaddressed AI Problems

Why Aren't Any AI Companies Watching Their Frontier Models to Make Sure They Don't G

17 AI Business Opportunities

Data centre firm Switch confidentially files for US IPO: Bloomberg

18 Copyright & Laws in AI Music

no new items today

19 General News

US offers \$1bn to Colombia on new right-wing president's first day of office

Remote AI Jobs

Worldwide-remote AI roles ranked against workflows/user_profile.md (n8n / Voiceflow / Relevance AI / Claude Code, entry-level). Senior/lead/principal and region-locked listings dropped. Sources: Remoteive, RemoteOK, WWR, Himalayas, HN, Green

TODAY'S SPOTLIGHT ROLE

Location Services Engineer | Maps Platform (Remote in Europe) - MapTiler

We Work Remotely . ★★★★★

WHY IT FITS

MapTiler is hiring: Location Services Engineer | Maps Platform (Remote in Europe). Headquarters: Zug, Switzerland URL: <https://www.maptiler.com/> F

[Apply →](#)

More roles today

WP dev (remote) - Snapscale ★★★★★

We Work Remotely

<https://weworkremotely.com/remote-jobs/snapscale-wp-dev-remote>

Web Designer - Steady Vision ★★★★★

We Work Remotely

<https://weworkremotely.com/remote-jobs/steady-vision-web-designer>

Website Builder / WordPress Specialist - Work and Study Travel ★★★★★

We Work Remotely

<https://weworkremotely.com/remote-jobs/work-and-study-travel-website-builder-wordpress-specialist>

AI Showcase Opportunities

Two blocks: Hackathons & Competitions (prize + deadline) and Accelerators & Incubators (who can apply, what you get, deadline). Every item has a direct apply link. India & worldwide.

HACKATHONS & COMPETITIONS

NEW · CUTC: Transform Hackathon

Prize: \$CAD 0

Deadline: August 8, 2026

https://cutc-transform.devpost.com/challenges/start_a_submission

Build with DataHub: The Agent Hackathon

Prize: \$20,500

Deadline: August 10, 2026

https://datahub.devpost.com/challenges/start_a_submission

AI Factory

Deadline: August 10, 2026

<https://lablab.ai/event/nativebuilder-build-without-limits>

ARC White-Box Estimation Challenge 2026

Deadline: August 10, 2026

<https://www.aicrowd.com/challenges/arc-white-box-estimation-challenge-2026>

Arm Create: AI Optimization Challenge

Prize: \$8,000

Deadline: August 14, 2026

https://arm-ai-optimization-challenge.devpost.com/challenges/start_a_submission

NTU InnovateX Hackathon 2026

Prize: \$20,000

Deadline: August 14, 2026

https://ntu-cctf-snz-innovatex-2026.devpost.com/challenges/start_a_submission

Canal Startup Sprint Hackathon

Prize: \$800

Deadline: August 14, 2026

https://canal-startup-sprint-hackathon.devpost.com/challenges/start_a_submission

ML Empowerment Build Challenge 2.0

Prize: \$1,000

Deadline: August 15, 2026

https://ml-empowerment-2.devpost.com/challenges/start_a_submission

02 . AI SHOWCASE OPPORTUNITIES (CONT.)

AceSAT Education AI-Agent

Prize: \$100

Deadline: August 15, 2026

https://acesat-ai-agent.devpost.com/challenges/start_a_submission

NEXORA Global Hackathon

Prize: \$0

Deadline: August 15, 2026

https://nexora-global-hackathon.devpost.com/challenges/start_a_submission

GatewayGS & The AEI Initiative: AI 4 Earth Hackathon

Prize: \$25

Deadline: August 15, 2026

https://gatewaygs-ai-4-earth-hackathon.devpost.com/challenges/start_a_submission

INNOVIK 6.0 : International Hackathon 2026

Prize: ₹ 200,000

Deadline: August 15, 2026

https://innovik-6-0-hackathon-2026.devpost.com/challenges/start_a_submission

Nexora

Prize: ₹ 30,000

Deadline: August 16, 2026

https://nexora.devpost.com/challenges/start_a_submission

Build with Gemini XPRIZE

Prize: \$2,000,000

Deadline: August 17, 2026

https://xprize.devpost.com/challenges/start_a_submission

YouCam API Skin AI & Apparel VTO Hackathon

Prize: \$6,000

Deadline: August 17, 2026

https://youcam-api.devpost.com/challenges/start_a_submission

ACCELERATORS & INCUBATORS

IndiaAI Startups Global - International Acceleration Programs for Indian AI Startups

Who can apply: Indian-registered AI startups (DPIIT-recognised).

What you get:

Placement into international acceleration tracks, plus government backing and global market access.

Deadline: Rolling - apply anytime

<https://indiaai.gov.in/>

Y Combinator (AI startups)

Who can apply: Any early-stage startup with a founding team; pre-launch and idea-stage are welcome.

What you get: ~\$500K on the standard deal, a 3-month program, and the lifelong YC network + Demo Day.

Deadline: Rolling - apply anytime

<https://www.ycombinator.com/apply>

Antler Residency (AI)

Who can apply: Solo founders welcome - no company or co-founder needed.

What you get: Pre-seed cash, a co-founder matching residency, and hands-on mentorship.

Deadline: Rolling - apply anytime

<https://www.antler.co/apply>

Techstars AI Accelerators

Who can apply: Early-stage startups; apply to the city/track that fits your team.

What you get: Seed investment, a 3-month mentor-driven program, and a global investor network.

Deadline: Rolling - apply anytime

<https://www.techstars.com/accelerators>

NVIDIA Inception

Who can apply: Any AI/ML startup at any stage - free to join.

What you get: GPU and cloud credits, go-to-market support, and introductions to VCs.

Deadline: Rolling - apply anytime

<https://www.nvidia.com/en-us/startups/>

Google for Startups AI Accelerator

Who can apply: Seed-to-Series-A AI startups; regional tracks including India.

What you get: An equity-free program, Google mentorship, and cloud credits.

Deadline: Rolling - apply anytime

<https://startup.google.com/programs/accelerator/>

Google Selects 20 AI Startups for 2026 Play Accelerator India Program - INDIA New England News

What you get:

Google Selects 20 AI Startups for 2026 Play Accelerator India Program - INDIA New England News

Deadline: Rolling - apply anytime

<https://news.google.com/rss/articles/CBMinwFBVV95cUxNR19sQVpGTvdhLXp5WkhNbWpHRTdpLUdCSDFJRjhXXzJJb0UwMkZ5bHIHe...>

Google Selects 20 AI Startups for 2026 Play Accelerator India Cohort - Siliconindia

What you get: Google Selects 20 AI Startups for 2026 Play Accelerator India Cohort - Siliconindia

Deadline: Rolling - apply anytime

<https://news.google.com/rss/articles/CBMizwFBVV95cUxQQzRhYUQxYTFuMTIaNmdFV1p0eFpOd0J4RDFOYjF3c19Ob0ZLTWd5ejE3YjF...>

YouTube Content Ideas

Three video ideas synthesised by the agent from the rest of this PDF, engineered to plausibly hit 10M views. Each ships with a hook, a thumbnail concept, and a 5-beat outline you can shoot the same week.

1. I Ran Two AI Agents That Talked To Each Other

HOOK (FIRST 8S)

Two terminals. Same laptop. One writes the code, the other tears it apart - and I never typed a word between them. This shipped last night and almost nobody has tried it yet.

Why it hits 10M: ['anthropic_claude_news - Claude Code sessions can now message each other and share context across terminals, so a real multi-agent setup needs zero framework, zero orchestration library and zero API keys. The entry barrier just collapsed and the tutorial gap is wide open.', 'ai_self_improvement_rsi - Time's 'Inside the Race to Make AI Build Itself' and Prime Intellect's open-sourced self-improving coding harness have pushed 'AI that improves itself' into mainstream anxiety. This video is the smallest honest version of that story a viewer can reproduce on their own machine tonight.', 'ai_search_trends - 'Claude Code: sessions can now message each other' is trending on Hacker News right now alongside 'what is claude' and 'claude ai app' rising on Google Trends in India, so search demand exists before the tutorials do.']

Thumbnail: Split screen of two terminal windows joined by a glowing arrow between them, one green and one red, with the words TALKING TO EACH OTHER stamped across the seam.

- Cold open: screen recording of session A messaging session B, B replying with a code review, and a test suite going green - all with hands off the keyboard.
- The setup in 90 seconds: two Claude Code sessions, the exact commands, what each one is told its job is.
- Watch it fail: the reviewer approves broken code, showing why an unsupervised agent pair drifts and needs a hard gate.
- The fix: give the reviewer a checklist and a test command, so verification is deterministic instead of vibes - the difference between a demo and something you'd ship.
- Where this goes: from August 14 Auto Mode is the Claude Code default, so this loop runs with far fewer approval prompts. Close on the honest limits and a link to the exact config.

2. The \$0.02 Agent That Cost Me \$400 In Electricity

HOOK (FIRST 8S)

Everyone tells you an AI agent costs pennies. A climate scientist metered his for eight weeks and found each one burns six hundred times the energy of a normal prompt. Here's what that actually did to my bill.

Why it hits 10M: ['anthropic_claude_news - Zeke Hausfather metered 3.2 billion tokens and roughly 170 kWh across eight weeks of Claude Code usage, about 600x a simple chat prompt per prompt. Concrete measured numbers are what turn a vague 'AI uses power' argument into a video people forward.', 'global_ai_news - AI datacenter air pollution and a town actively lobbying to host a datacenter are both in the news cycle this week, so the personal cost story sits inside a live public argument instead of floating alone.', 'ai_business_automation - the professional-services piece shows firms rebuilding delivery around AI while billing by the hour; if agent runs have a real energy and token cost, that pricing model is where the number lands. This gives the video a business punchline, not just an outrage one.']

Thumbnail: An electricity meter dial spinning into a blur, with a laptop terminal reflected in its glass and a hand-drawn 600x circled in red.

- Cold open: the metered number on screen - 3.2 billion tokens, 170 kWh - then the question, is my agent worth this.
- Where the tokens actually go: instrument one agent run and show the loop re-reading files, retrying tools, and quietly 10x-ing the visible prompt.

3. Nobody Is Watching The AI. I Built The Watcher.

HOOK (FIRST 8S)

Not one major AI lab continuously monitors what its models do after release. I gave an agent real credentials, pointed a second one at it as a tripwire, and recorded everything that happened next.

Why it hits 10M: ['unaddressed_ai_problems - the frontier-model monitoring gap is a named, still-open problem: labs ship agents with real permissions and publish no live behavioural monitoring. A video that builds the missing thing is far stronger than one that complains about it.', 'unaddressed_ai_problems - the OpenAI/Hugging Face incident timeline shows an internal misconfiguration cascading into an outage of shared infrastructure, which is exactly the failure class an audit-and-alert layer catches. It gives the build a real, recent, verifiable motivation.', 'new_ai_tools - Backflip AI turning 3D scans into editable CAD is the kind of agent that gets handed real filesystem and account access with no oversight layer, so the audience for a watcher is already there and growing.']

Thumbnail: A dark dashboard of agent actions scrolling past with one row highlighted blood-red and an eye icon watching it, captioned NOBODY IS WATCHING.

- Cold open: the alert firing when the monitored agent tries something it should not, before any explanation of what was built.
- The gap: what labs actually publish about deployed model behaviour versus what an enterprise handing over credentials assumes is happening.
- Build it live: log every agent tool call to a structured store, define what abnormal looks like, and alert on it - under an hour, no framework.
- Red-team my own watcher: three ways the monitored agent slips past it, and the two fixes that hold.
- Why this is a portfolio piece: it is the audit trail every compliance-bound client will ask for, and almost nobody applying for these jobs has built one.

Viral AI on YouTube

Verified-viral picks from the last 7 days (2 long: Global + India, 1 Global Short). Floors: long >= 100K, short >= 500K. No view counts. No automation angle.

OpenAI's Genius Comeback : How Altman Turned a \$27B Loss Into an EMPIRE? - Think School

video . Global Automation

WHY IT WENT VIRAL

OpenAI's Genius Comeback : How Altman Turned a \$27B Loss Into an EMPIRE?.

https://www.youtube.com/watch?v=opLMb_oYfMI

A compilation of chaotic situations taken to the extreme #ai #shorts #chunkychaos Dream 10K Sub - Chunky Chaos

short . Global - Short

WHY IT WENT VIRAL

A compilation of chaotic situations taken to the extreme #ai #shorts #chunkychaos Dream 10K Sub.

<https://www.youtube.com/watch?v=AYaFeelILDQ>

video . India Automation

WHY IT WENT VIRAL

.

Viral Instagram Reels

Top-engagement AI reels in the last 24h - one Global, one India. Ranked by likes + comments, AI-filtered, URL-checked. Reverse-engineer the hook for your own reels.

Global

WHY IT WORKS

@ - 0 likes / 0 comments in 24h. #. Automation angle: Study the first-2s hook + caption format - replicate the structure for your AI-automation Instagram reels.

India

WHY IT WORKS

@ - 0 likes / 0 comments in 24h. #. Automation angle: Study the first-2s hook + caption format - replicate the structure for your AI-automation Instagram reels.

Quantum + AI

Stories addressing both quantum (qubits, quantum hardware/algorithms) and AI/ML. Pure-quantum or pure-AI items are dropped.

QuanTiMedAI: Quantum-Enhanced Time-Series Model guided by Agentic AI for Cardiac Arrest Mortality Prediction

QUICK TAKE

QuanTiMedAI pairs a quantum-enhanced time-series model with an agentic AI controller to predict ICU mortality after cardiac arrest, using the agent to select and orchestrate the modelling steps rather than fixing them by hand. It is a clean example of the hybrid pattern: a quantum-enhanced component inside a pipeline an LLM agent drives. Automation angle: Note the architecture, not the physics - an agent choosing which model to run per case is a pattern you can rebuild classically today.

arXiv AI . 2026-08-06 . ★★★★★
<https://arxiv.org/abs/2608.06294v1>

Neural Flux Attachment: From Bose Condensates to Chiral Topological Matter

QUICK TAKE

ChernFormer combines a fermionic transformer with a fixed Chern-Simons phase so that one neural wave function can describe both a Bose condensate and a chiral topological liquid. It is neural-network quantum states applied to a regime where those two phases were previously modelled separately. Automation angle: Transformers are quietly replacing bespoke solvers in physics - a reminder that domain-specific fine-tuning, not raw model scale, is where applied AI work actually lives.

arXiv Quantum Physics . 2026-08-06 . ★★★★★
<https://arxiv.org/abs/2608.06168v1>

Argonne's ChemGraph opens quantum chemistry to non-experts, tracks AI model performance - R&D World

QUICK TAKE

Argonne's ChemGraph puts an agentic interface over quantum-chemistry workflows so non-specialists can run simulations, and it tracks how well the underlying AI models perform on those tasks. The built-in performance tracking is the notable part - the tool measures its own reliability. Automation angle: Copy the pattern: ship an eval dashboard alongside any agent you hand a client. Self-reported accuracy is what turns a demo into something someone will pay for.

Quantum + AI (via Google News) . 2026-08-06 . ★★★★★
<https://news.google.com/rss/articles/CBMitwFBVV95cUxPUUhSZjJYbDVFTHV3dkNXR1Mts2lhVjg2dVJqcE1tSkV0dzc1ckZBdTVaakZ4d...>

AI Models Run 93% Faster On Qualcomm Chips, Multiverse Computing Says - Quantum Zeitgeist

QUICK TAKE

Multiverse Computing says its quantum-inspired compression makes AI models run about 93% faster on Qualcomm chips - classical tensor-network methods borrowed from quantum simulation, not quantum hardware. It targets on-device inference, where the memory wall bites hardest. Automation angle: Quantum-inspired compression is usable today with no QPU - worth watching if you want agents running locally instead of per-token billed.

Quantum + AI (via Google News) . 2026-08-06 . ★★★★★

<https://news.google.com/rss/articles/CBMid0FVX3lxTFBGU3lxX3lwUjBRd1FVdWNLOTR0TDNScVpqeFdmWU91aU5lNzRQTmJnN0lwWn...>

Can quantum take the load off AI's power problem? - InformationWeek

QUICK TAKE

InformationWeek examines whether quantum computing can relieve AI's power problem, and the honest answer in the reporting is: not soon, and not for general training workloads. The realistic near-term overlap is narrow optimisation problems, not model training. Automation angle: Useful counter-hype material - being the person who can say why quantum will not fix inference costs this decade is worth more than repeating the pitch.

Quantum + AI (via Google News) . 2026-08-07 . ★★★★★

<https://news.google.com/rss/articles/CBMiowFBVV95cUxPaGM2YW55QnJfTk5CODBockJDMDQtSmc0ZWZXMjNfcjNTbW9iT2xueIJRML...>

Gartner Predicts Enterprise AI Workloads at Scale Will Not Run on Quantum Hardware Through 2028 - Scientific Computing World

QUICK TAKE

Gartner forecasts that enterprise AI workloads at scale will not run on quantum hardware through 2028, framing today's quantum-plus-AI announcements as research rather than procurement decisions. It gives buyers a defensible date to plan against. Automation angle: Quote the date when a client asks whether to wait for quantum: the answer is build classical now, revisit after 2028.

Quantum + AI (via Google News) . 2026-08-05 . ★★★★★

<https://news.google.com/rss/articles/CBMizwFBVV95cUxPU3lPS0wyUEVDeWZEEHBzc2hGLVFuMUlvQ2pGeVhqb2d6OU8wVTIKLUZIT...>

AI Self-Improvement (RSI)

Recursive self-improvement, AGI progress, alignment research.

Inside the Race to Make AI Build Itself - Time Magazine

QUICK TAKE

Time's feature goes inside the labs explicitly trying to get AI systems to build their next generation themselves, treating automated AI research as an engineering programme rather than a thought experiment. It is the clearest mainstream framing yet of self-improvement as a funded roadmap. Automation angle: The skill that transfers is closing the loop - an agent that runs an experiment, scores it, and edits its own next attempt is buildable at small scale and demos extremely well.

AI Self-Improvement (via Google News) . 2026-08-07 . ★★★★★

<https://news.google.com/rss/articles/CBMiwwFBVV95cUxQYWVQUnNEbTExb2lFMkQ3Yzh3cHIYQTNyYzYzLXluVU9NbUQyNmhT...>

Prime Intellect Open Sources Prime Agent Self-Improving AI Coding Harness - Open Source For You

QUICK TAKE

Prime Intellect open-sourced Prime Agent, a self-improving coding harness that revises its own scaffolding across attempts instead of relying on a fixed prompt and toolset. Because the harness is public, the improvement loop can be inspected rather than taken on faith. Automation angle: Read this repo - an open self-improving harness is the cheapest way to learn agent-loop design without paying for frontier experiments.

AI Self-Improvement (via Google News) . 2026-08-07 . ★★★★★

<https://news.google.com/rss/articles/CBMidkFVX3lxTFBJRFN2V3p6OVgwMFhUdHdja1VFeUhhKVHR4a0dZLUIMNGw0Z3J6ZIM3dkhxODh...>

Exclusive: Mirendil inks \$100M+ Google Cloud deal to scale self-improving AI - TechCrunch

QUICK TAKE

AI lab Mirendil signed a Google Cloud deal worth over \$100 million specifically to scale self-improving AI systems, per TechCrunch. Compute commitments of that size indicate self-improvement work has moved from papers into capacity planning. Automation angle: Follow where the compute contracts go - they lead the hiring wave by roughly a year, which is useful when picking what to specialise in.

AI Self-Improvement (via Google News) . 2026-08-06 . ★★★★★

<https://news.google.com/rss/articles/CBMirAFBVV95cUxON0xsdzZxSURyVINfMFJKdV9zYjcweGJ0THJvVmFtQnlxWXp0YjdUFZFNrkRk...>

On-Policy Self-Distillation without Any Supervision

QUICK TAKE

This paper strips external supervision out of on-policy self-distillation for LLMs entirely: no ground-truth labels, no external verifier, the model improves from its own on-policy behaviour. Removing the supervision requirement is what makes an AI training loop genuinely self-sustaining. Automation angle: Self-consistency scoring is the practical takeaway - use the same trick to grade LLM agent outputs when you have no labelled data.

arXiv Machine Learning . 2026-08-06 . ★★★★★

<https://arxiv.org/abs/2608.06296v1>

DASH: Divergence-Adaptive Supervision Horizons for On-Policy Self-Distillation of Reasoning Models

QUICK TAKE

DASH adapts the supervision horizon during on-policy self-distillation of reasoning models, addressing the sparse sequence-level signal that limits reinforcement learning with verifiable rewards. Adaptive horizons give the model denser feedback about which part of its own reasoning to fix. Automation angle: Same lesson at the application layer: score intermediate agent steps, not just the final answer, or you never learn where it went wrong.

arXiv AI . 2026-08-06 . ★★★★★

<https://arxiv.org/abs/2608.06243v1>

Future AGI: Open-source platform for shipping self-improving AI agents - Help Net Security

QUICK TAKE

Future AGI released an open-source platform for shipping self-improving AI agents in production, packaging the evaluate-then-adapt loop as deployable infrastructure rather than a research artefact. Productionising that loop is the hard part most teams skip. Automation angle: An open self-improvement platform is a fast route to a credible portfolio project - deploy one agent on it and document the loop.

AI Self-Improvement (via Google News) . 2026-08-05 . ★★★★★

<https://news.google.com/rss/articles/CBMirgFBVV95cUxOTTYtZU5vdUZyVi1xX1Q5cnhwZ0JZYXBaV2dZaWR2eVQ3Zk9zejc2c1RmN0die...>

Hitting and Breaking Limits to Reach Unlimited Recursively Self Improving AI - NextBigFuture.com

QUICK TAKE

NextBigFuture walks through the physical and economic ceilings - compute, energy, data - that would cap recursive self-improvement, and what breaking each would require. It is a useful counterweight to unbounded takeoff framings. Automation angle: Constraint-first thinking is a good interview habit: name the bottleneck before the capability when asked where AI is going.

AI Self-Improvement (via Google News) . 2026-08-06 . ★★★★★

<https://news.google.com/rss/articles/CBMijwFBVV95cUxNaWlkTzJnYmVkcGV2bGxsLVFFUkNJa21nbVpyeEdSM2t1MUUs2OGFuYm1XSD...>

Elon Musk's AI Vision

xAI & Grok news, Elon Musk's AI views, statements & predictions.

Grok Imagine 2.0: xAI's image generator is now second only to OpenAI - tbreak.com

QUICK TAKE

xAI shipped Grok Imagine 2.0, its second-generation image generator, and it now sits second only to OpenAI on public image leaderboards. For xAI this is the first time a Grok model has ranked at the frontier of a category rather than trailing it. Automation angle: Image generation is now a two-horse race - if your portfolio has a visual-content pipeline, wire it behind an adapter so you can swap Grok Imagine and GPT-Image without touching the workflow.

xAI News (via Google News) . 2026-08-08 . ★★★★★

<https://news.google.com/rss/articles/CBMiZ0FVX3IxtFBIZGtZNEk2UmpfTVFyVHltUIJnVGFjbXVtQnpwZVF2LU42Sy10OXNBUGtZc2lvdHh...>

xAI's Imagine Image 2.0 lands just behind OpenAI's GPT-Image-2 in Arena benchmarks - the-decoder.com

QUICK TAKE

The Decoder's benchmark write-up puts xAI's Imagine Image 2.0 just behind OpenAI's GPT-Image-2 in Arena scoring, alongside new editing tools like Magic Wand and multi-reference conditioning. Editing controls, not raw fidelity, are where the two models now differentiate. Automation angle: Multi-reference conditioning is what makes brand-consistent asset generation possible - a concrete, sellable automation for small e-commerce and content clients.

xAI News (via Google News) . 2026-08-08 . ★★★★★

<https://news.google.com/rss/articles/CBMiqAFBVV95cUxNRDkzMDdVYIB3NIRvR2dYenhhcniPVEd5NzdLczdfc0ZzLXVJVkJsZExpZ2hFe...>

Grok Imagine 2.0 Pushes xAI Deeper Into OpenAI and Midjourney Territory - Startup Fortune

QUICK TAKE

Coverage frames Grok Imagine 2.0 as xAI pushing directly into territory held by OpenAI and Midjourney, rather than serving Grok users as a side feature. It is consistent with Musk's stated position that xAI should compete across every modality instead of specialising. Automation angle: Watch pricing here - when a third serious player enters image generation, per-image cost usually falls, which changes the margins on any content-automation service you sell.

xAI News (via Google News) . 2026-08-08 . ★★★★★

<https://news.google.com/rss/articles/CBMinwFBVV95cUxOdIgtbDJKT1dqWi1BZV9KamVfWWZsc3Z1X2dwQVRqdDBsX0t0eWUzUmU5V...>

SpaceX's \$60 Billion Cursor Acquisition Could Close Next Week; New Products May Adopt Grok Branding - finance.biggo.com

QUICK TAKE

Reports say SpaceX's roughly \$60 billion acquisition of Cursor could close next week, with new products possibly moving under Grok branding. If it lands, Musk would own a frontier coding agent outright - the clearest signal yet that he views developer tooling, not chat, as the strategic layer of AI. Automation angle: Keep your coding workflow portable across Cursor / Claude Code / Codex; an ownership change of this size usually resets pricing and terms.

xAI News (via Google News) . 2026-08-08 . ★★★★★

<https://news.google.com/rss/articles/CBMidkFVX3lxTE14TDFWdTdqNkluYktOTlp4NmNKVVplU3pZd2VfU1dJVjREaEw1OGtqSGthcmRwe...>

Imagine Image 2.0 - X.ai

QUICK TAKE

xAI's own release post for Imagine Image 2.0 is the primary source behind this week's Grok image-model coverage. It is worth reading directly for the model card and feature list rather than relying on the benchmark write-ups. Automation angle: Read vendor release notes before the news cycle summarises them - quoting model cards accurately is a cheap credibility win in interviews.

xAI News (via Google News) . 2026-08-08 . ★★★★★

<https://news.google.com/rss/articles/CBMiT0FVX3lxTE9HTVhhdDlteUJBZHduTWhqS0U5MW1NaFVPekdTZjNBbUd3VnNqaFlnd0FNyVI...>

AI Model Benchmarks

Top model per category (Text/LLM, Coding, Image, Video, Music, Audio) as a table, then benchmark news.

Task	Best	Runner-up	Benchmark
Text / LLM	Claude Opus 4.x	GPT-5.x	LMSYS Chatbot Arena (blended Elo)
Coding	Claude Fable 5 (via Claude Code)	GPT-5.6 Sol Ultra (via Codex)	SWE-bench Verified + one-shot build tests
Image Generation	OpenAI GPT-Image-2	xAI Imagine Image 2.0	LMarena image leaderboard (Aug 2026)
Video Generation	Sora 2	Google Veo 3	VBench 2.0 + human-eval
Music / Audio	Suno v5	Udio v2.5	MusicArena
Reasoning / Math	GPT-5.x (extended thinking)	Claude Opus 4.x	AIME / FrontierMath

Learning When to Trust via Selective Context Preference Optimization

QUICK TAKE

Selective Context Preference Optimization tackles a specific failure mode: models trained to resist misleading external context also learn to ignore context that was correct. The method teaches selective trust instead of blanket resistance. Automation angle: Directly relevant to RAG - measure how often your retriever's context makes the answer worse, not just how often it helps.

arXiv AI . 2026-08-06 . ★★★★★
<https://arxiv.org/abs/2608.06377v1>

The Low Frequency Trap: Video Language Models Fail at Simple Event Bookkeeping

QUICK TAKE

This benchmark isolates event counting, rate and duration in video and finds video language models fail at simple event bookkeeping once those variables are separated from visual complexity. Standard video benchmarks entangle the factors and hide the failure. Automation angle: Before selling any video-understanding automation, test counting explicitly - it is the failure your client will notice first.

arXiv AI . 2026-08-06 . ★★★★★
<https://arxiv.org/abs/2608.06361v1>

New AI Tools

Latest AI tools with cost & feature notes - biased to Claude Code / n8n / Voiceflow / Relevance AI / MCP / agent builders.

Backflip AI turns 3D scans into editable CAD models in minutes instead of hours

QUICK TAKE

Backflip AI released a model that turns raw 3D scans into fully editable parametric CAD, collapsing a task that normally takes hours of expert work. CEO Greg Mark's pitch is that most scan-to-CAD work today is manual retracing. Automation angle: Scan-to-CAD is a textbook niche automation - fabrication and furniture shops will pay for a working intake pipeline long before they pay for a chatbot.

The Decoder . 2026-08-08 . ★★★★★

<https://the-decoder.com/backflip-ai-turns-3d-scans-into-editable-cad-models-in-minutes-instead-of-hours/>

Apple says Mac users in China can connect to Alibaba's Qwen AI service

QUICK TAKE

Apple published a guide letting Mac users in China connect Siri and Writing Tools to Alibaba's Qwen, its first sanctioned third-party model routing on the desktop. It is a regulatory workaround, but it also proves the OS-level assistant is becoming a swappable backend. Automation angle: Design your agents to talk to a model router, not a vendor SDK - the platforms themselves are now normalising model substitution.

Economic Times Tech . 2026-08-08 . ★★★★★

<https://economictimes.indiatimes.com/tech/technology/apple-says-mac-users-in-china-can-connect-to-alibabas-qwen-ai-service/article...>

Indian AI Industry

India-specific AI news only - Indian AI startups, AI policy, AI products.

Government examines Meta's recommendation system for intermediary status

QUICK TAKE

India's IT ministry is examining whether Meta's AI recommendation systems and paid content promotion still qualify for safe-harbour intermediary status under the IT Act, citing AI deepfakes and content moderation as the triggers. If a ranking model counts as publishing, every AI recommender operating in India inherits publisher liability. Automation angle: Any AI ranking or recommendation feature you ship for an Indian client now needs an audit trail of why an item was surfaced - log it from day one, it is far harder to retrofit.

YourStory . 2026-08-08 . ★★★★★

<https://yourstory.com/2026/08/govt-examines-meta-recommendation-system-intermediary-status>

Anthropic & Claude Code

Anthropic company news + Claude Code feature updates (new commands, hooks, MCP, slash commands, agents, model rollouts).

Anthropic sets Claude Code to Auto Mode by default to protect developers from bad approvals

QUICK TAKE

From August 14 Anthropic makes Auto Mode the default permission mode in Claude Code for Pro, Max and Team plans, so a classifier - not a wall of yes/no prompts - decides which commands need review. Anthropic says the classifier caught 89% of dangerous commands in internal tests, and the change is aimed squarely at approval fatigue, where developers rubber-stamp every prompt and stop reading them. Automation angle: Re-test your Claude Code hooks and deny-lists before the 14th - any automation that assumed a manual approval gate now runs unattended.

The Decoder . 2026-08-08 . ★★★★★

<https://the-decoder.com/anthropic-sets-claude-code-to-auto-mode-by-default-to-protect-developers-from-bad-approvals/>

Claude Code sessions can now talk to each other and share context across terminals

QUICK TAKE

Claude Code sessions running in parallel on macOS and Linux can now message each other, share context and query each other's status. That turns several terminals into a crude multi-agent setup without any orchestration framework on top. Automation angle: Build a two-session pattern - one session writing code, one reviewing and running tests - and record it; it is a portfolio demo of multi-agent orchestration that costs nothing but a second terminal.

The Decoder . 2026-08-08 . ★★★★★

<https://the-decoder.com/claude-code-sessions-can-now-talk-to-each-other-and-share-context-across-terminals/>

AI agents use roughly 600 times more energy than a simple chat prompt

QUICK TAKE

Climate scientist Zeke Hausfather metered eight weeks of his own Claude Code usage: 3.2 billion tokens and roughly 170 kWh of datacenter electricity, or about 600x the energy of a single chat prompt. The gap is the agentic loop itself - tool calls, re-reads and retries multiply token spend far beyond what the user types. Automation angle: Add a token/cost budget line to every agent you ship; clients increasingly ask what an agent run costs, and 'I measured it' is a differentiator when you have no formal experience.

The Decoder . 2026-08-08 . ★★★★★

<https://the-decoder.com/ai-agents-use-roughly-600-times-more-energy-than-a-simple-chat-prompt/>

Apple and OpenAI escalate legal battle over devices

QUICK TAKE

Apple and OpenAI's device dispute escalated further this week, landing while OpenAI is reportedly valued near \$852 billion and weighing an IPO. The fight matters for anyone building on assistant platforms, because it decides which models get first-party placement on the hardware people actually carry. Automation angle: Do not hard-wire an agent to one vendor's assistant surface - keep the model layer swappable so a platform fight never strands your build.

Economic Times Tech . 2026-08-08 . ★★★★★

<https://economictimes.indiatimes.com/tech/technology/apple-and-openai-escalate-legal-battle-over-devices/articleshow/133046522.cms>

AI Automation & Business

How AI automation is changing work across industries: real deployments, ROI, sector updates.

AI is reimagining service companies; Ola Electric narrows loss in Q1

QUICK TAKE

Professional-services firms are rebuilding their delivery model around AI rather than bolting it on: Northbound Advisors is training a proprietary tool on eight years of diligence, CFO advisory, legal, tax and risk engagements. The economics of consulting - bill by the hour, staff by the headcount - break when the firm's own back catalogue becomes the product. Automation angle: The moat is the archive, not the model. Offer small firms a 'turn your past deliverables into a retrieval agent' engagement; it is buildable with n8n plus a vector store and needs no proprietary research.

YourStory . 2026-08-08 . ★★★★★

<https://yourstory.com/2026/08/ai-is-reimagining-service-companies>

Readers rate AI-generated short stories higher than human ones until they learn a machine wrote them

QUICK TAKE

In a study of more than 2,500 readers, people could not distinguish ChatGPT-written short stories from human ones - they performed at chance, and actually rated the AI texts higher until they were told a machine wrote them. The penalty is attached to disclosure, not to quality. Automation angle: For content clients this is the whole compliance question in one study: build disclosure into the workflow deliberately, because the reputational cost lands when the label appears, not when the text does.

The Decoder . 2026-08-08 . ★★★★★

<https://the-decoder.com/readers-rate-ai-generated-short-stories-higher-than-human-ones-until-they-learn-a-machine-wrote-them/>

Global AI News

Worldwide AI news only (US, EU, China, Japan, Korea, Middle East) - no general tech, no non-AI business news.

AI Data Centers Are Causing Unfathomable Amounts of Air Pollution, and It Gets Worse With Each New One They Build

QUICK TAKE

New reporting puts hard numbers on the air-quality cost of AI datacenters, much of it from the on-site gas generators used to cover grid shortfalls, and the footprint grows with every new build. The industry shows no sign of slowing the buildout. Automation angle: Efficiency is becoming a sales argument, not just a cost line - quantify tokens saved per run in your demos, because buyers are starting to be asked about it.

Futurism . 2026-08-08 . ★★★★★

<https://futurism.com/future-society/ai-data-centers-air-pollution-generators-worse>

Fields Medalist who published a paper on AI-driven human extinction now works for OpenAI

QUICK TAKE

Newly minted Fields Medalist Jacob Tsimmerman is leaving the University of Toronto for OpenAI to work on safety, having recently published an analysis of scenarios where AI contributes to human extinction. The hire is a data point on where mathematical talent is now flowing. Automation angle: Safety and evaluation work is one of the few AI specialisms still hiring people without a research pedigree - building a rigorous eval harness is a legitimate portfolio piece.

The Decoder . 2026-08-08 . ★★★★★

<https://the-decoder.com/fields-medalist-who-published-a-paper-on-ai-driven-human-extinction-now-works-for-openai/>

Weird Town Beggars for AI Data Center

QUICK TAKE

A town is actively lobbying to host an AI datacenter, inverting the usual pattern of local opposition, and the piece is a caution about what communities trade for the tax base. Datacenter siting is becoming local politics everywhere. Automation angle: Compute geography drives latency and price - know which region your client's data can legally sit in before you pick an inference provider.

Futurism . 2026-08-08 . ★★★★★

<https://futurism.com/artificial-intelligence/weird-town-begs-ai-data-center>

How to Disable Gemini in Gmail and Google Docs

QUICK TAKE

Google has pushed Gemini toolbars and prompts into Docs and Gmail, and Wired walked through turning them off. The interesting part is that opting out is now a documented user need for a default-on assistant. Automation angle: Default-on AI in office suites is your competition for simple tasks - pitch automations that span tools (inbox to CRM to sheet), which the in-app assistants still cannot do.

Wired AI . 2026-08-08 . ★★★★★

<https://www.wired.com/story/how-to-disable-the-gemini-ai-features-in-gmail-and-google-docs/>

OpenAI flags its new Astra model as potentially reaching the highest cybersecurity risk level for the first time

QUICK TAKE

OpenAI says internal testing of its in-development Astra model showed cybersecurity capability strong enough that it can no longer rule out the highest risk tier in its own preparedness framework - the first time it has said this about a model. Development activity around Astra has been paused as a result. Automation angle: Expect tighter gating on agentic capabilities across all providers; build with least-privilege credentials now so a capability lockdown does not break your pipeline later.

The Decoder . 2026-08-08 . ★★☆☆☆

<https://the-decoder.com/openai-flags-its-new-astra-model-as-potentially-reaching-the-highest-cybersecurity-risk-level-for-the-first-ti...>

What People Are Searching For

Hottest AI topics right now: Google Trends rising queries (global + India), Hacker News top AI stories, and Reddit AI subreddits (last 24h).

generative ai meaning

QUICK TAKE

Trending across: google_trends_rising, google_trends_top. Geo: india.

★★★★★

<https://trends.google.com/trends/explore?q=generative+ai+meaning&date=now+1-d&geo=IN>

open generative ai

QUICK TAKE

Trending across: google_trends_rising, google_trends_top. Geo: india.

★★★★★

<https://trends.google.com/trends/explore?q=open+generative+ai&date=now+1-d&geo=IN>

Oracle bans AI-generated code from OpenJDK

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://app.dealroom.co/news/feed/oracle-bans-ai-generated-code-from-openjdk-despite-ellison-s-claim-oracle-isn-t-writing-its-own...>

what is claude

QUICK TAKE

Trending across: google_trends_top. Geo: india.

★★★★★

<https://trends.google.com/trends/explore?q=ai+claude&date=now+1-d&geo=IN>

Timeline of the OpenAI accidental attack against Hugging Face

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://simonwillison.net/2026/Aug/7/openai-timeline/>

Managing AI Coding Costs at Scale

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://www.databricks.com/blog/managing-ai-coding-costs-scale>

Gentoo bugzilla closed due AI bot scraper overload

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://social.treehouse.systems/@mgorny/117058483039362779>

Lost my phone at the office. Claude suggested tracking Bluetooth signal strength

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://twitter.com/un1c0rnioz/status/2084686552299634805>

what is generative ai

QUICK TAKE

Trending across: google_trends_top. Geo: india.

★★★★★

<https://trends.google.com/trends/explore?q=what+is+generative+ai&date=now+1-d&geo=IN>

Mythos social engineering AISI INC-2026-07-28-01

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://web.archive.org/web/20260731053721/http://github.com/ancaferro/myNetwork/pull/3>

generative ai course

QUICK TAKE

Trending across: google_trends_top. Geo: india.

★★★★★

<https://trends.google.com/trends/explore?q=generative+ai+course&date=now+1-d&geo=IN>

Should AI labs be treated like the owners of dangerous animals?

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://www.economist.com/science-and-technology/2026/08/06/should-ai-labs-be-treated-like-the-owners-of-dangerous-animals>

The Claudyssey: A line-for-line translation of Homer's Odyssey by Claude Fable 5

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://theclaudyssey.com/>

perplexity ai

QUICK TAKE

Trending across: google_trends_rising. Geo: india.

★★★★★

<https://trends.google.com/trends/explore?q=perplexity+ai&date=now+1-d&geo=IN>

Making an AI bid writer refuse to lie

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://ailucius.com/blog/making-an-ai-bid-writer-refuse-to-lie>

claude ai app

QUICK TAKE

Trending across: google_trends_rising. Geo: india.

★★★★★

<https://trends.google.com/trends/explore?q=claude+ai+app&date=now+1-d&geo=IN>

Claude Code: Starting August 14, auto mode will be the default permission mode

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://twitter.com/ClaudeDevs/status/2085794862608318627>

YouTube's AI Detection Kicked Us in the Face

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

https://twitter.com/Kurz_Gesagt/status/2083191397981561232

Claude Code: sessions can now message each other

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://twitter.com/ClaudeDevs/status/2085817074816070014>

Cloudflare launches Kitesurf, a browser built for AI agents

QUICK TAKE

Trending across: hn. Geo: global.

★★★★★

<https://techcrunch.com/2026/08/07/cloudflare-launches-kitesurf-a-browser-built-for-ai-agents/>

Unaddressed AI Problems

Real problems in AI that nobody is solving - gaps & unmet needs.

Why Aren't Any AI Companies Watching Their Frontier Models to Make Sure They Don't Go on Hacking Sprees?

QUICK TAKE

Problem: nobody is continuously monitoring deployed frontier models for autonomous misuse. The piece argues runtime containment is far more tractable than labs imply, yet no major lab publishes live behavioural monitoring of models already in production. Who it hits: every enterprise handing an agent real credentials, with no vendor-side tripwire behind it. Why it is still open: monitoring costs inference margin and surfaces incidents labs would rather not disclose. Automation angle: This is a buildable gap - an agent-action audit log with anomaly alerts is a genuinely sellable product and a strong portfolio project.

Futurism . 2026-08-08 . ★★★★★

<https://futurism.com/future-society/ai-companies-watching-frontier-models-hacking-sprees>

Now we have a timeline of the OpenAI accidental attack against Hugging Face

QUICK TAKE

Problem: AI crawlers can take down critical open infrastructure by accident, and nothing prevents it. OpenAI's Black Hat timeline of the Hugging Face incident shows an internal misconfiguration escalating into an outage of a dependency the whole ecosystem builds on. Who it hits: every maintainer of a free model or dataset host, and everyone downstream of them. Why it is still open: there is no shared rate-limit contract or crawler identity standard between labs and the hosts they depend on. Automation angle: Pin and mirror the model weights and datasets your automations pull at runtime - an upstream outage should not be able to stop your client's workflow.

Simon Willison's Blog . 2026-08-08 . ★★★★★

<https://simonwillison.net/2026/Aug/8/now-we-have-a-timeline-of-the-openai-accidental-attack-against-h/#atom-everything>

AI Business Opportunities

Emerging business opportunities in AI - India & world.

Data centre firm Switch confidentially files for US IPO: Bloomberg

QUICK TAKE

Switch has confidentially filed for a US IPO on the back of large-scale campuses supplying the power, cooling and connectivity that GPU clusters need. Public-market appetite for AI infrastructure is a proxy for how long the compute buildout is expected to run. What to build: the unglamorous operations layer around these sites - capacity reporting, vendor-invoice reconciliation, incident summarisation - sold to colo and MSP operators who are scaling headcount to keep up. Automation angle: Infrastructure booms create back-office pain first; that back office is exactly what an n8n-plus-LLM workflow automates well.

Economic Times Tech . 2026-08-08 . ★★★★★

<https://economictimes.indiatimes.com/tech/technology/data-centre-firm-switch-confidentially-files-for-us-ipo-bloomberg/articleshow/13...>

SK Hynix board approves \$38 billion investments for South Korea's Yongin, Cheongju chip plants

QUICK TAKE

SK Hynix's board approved roughly \$38 billion for the Yongin and Cheongju fabs, with Yongin's Y2 fab dedicated to DRAM - the memory that gates AI training and inference throughput. Memory, not just GPUs, is now the constraint being capitalised. What to build: for whom this matters is mid-size hardware distributors and electronics importers - component price and lead-time monitoring, quoted automatically to their customers, is a concrete first automation product. Automation angle: Why now: supply-chain volatility plus cheap structured extraction makes price/lead-time intelligence sellable to firms still tracking it in spreadsheets.

Economic Times Tech . 2026-08-08 . ★★★★★

<https://economictimes.indiatimes.com/tech/technology/sk-hynix-board-approves-38-billion-investments-for-south-koreas-yongin-che...>

Copyright & Laws in AI Music

AI music copyright lawsuits, regulations, fair-use rulings, licensing - India & world.

No new items in the last 24h.

General News

Top world & India headlines outside of AI.

US offers \$1bn to Colombia on new right-wing president's first day of office

QUICK TAKE

Washington offered Colombia \$1 billion in aid on the first day of right-wing president Abelardo de la Espriella's term, who used his maiden speech to promise an 'all-out war' on what he called narco-terrorism.

BBC World . 2026-08-08 . ★★★★★

https://www.bbc.co.uk/news/articles/cy9wy3y0e5wo?at_medium=RSS&at_campaign=rss

Trump's ex-lawyer Todd Blanche narrowly confirmed as US attorney general

QUICK TAKE

Todd Blanche, Donald Trump's former criminal defence lawyer, was narrowly confirmed as US attorney general despite unusual pushback from some Republican senators, placing a Trump personal attorney atop federal law enforcement.

BBC World . 2026-08-08 . ★★★★★

https://www.bbc.co.uk/news/articles/cewr898jy8go?at_medium=RSS&at_campaign=rss

Child among three killed in Russian missile attacks near Kyiv

QUICK TAKE

Russian missile strikes near Kyiv killed three people including a child, continuing overnight attacks as President Zelensky warned Ukraine's stock of interceptor missiles is running down.

BBC World . 2026-08-08 . ★★★★★

https://www.bbc.co.uk/news/articles/cz7dy8gq99eo?at_medium=RSS&at_campaign=rss

Typhoon Dolphin batters Japan's Okinawa before bearing down on China

QUICK TAKE

Typhoon Dolphin battered Japan's Okinawa, cutting power to at least 44,000 buildings and injuring five people, and is now bearing down on the Chinese coast.

BBC World . 2026-08-08 . ★★★★★

https://www.bbc.co.uk/news/articles/cvglp2zlvrl0?at_medium=RSS&at_campaign=rss

Spain imposes border controls against Italy as row over Ceuta migrant influx intensifies

QUICK TAKE

Spain imposed border controls against Italy as the row over roughly 78,000 migrants crossing from Morocco into the Spanish exclave of Ceuta escalated, with Italy having introduced controls of its own.

BBC World . 2026-08-08 . ★★★★★

https://www.bbc.co.uk/news/articles/c5yv5l6yr5ko?at_medium=RSS&at_campaign=rss

Iran Claims Hormuz Deal Close, But Adds Caveat For Opening Strait

QUICK TAKE

Iran said a deal on the Strait of Hormuz is close but attached fresh conditions to reopening it; a US official said Washington expects an agreement soon between Iran and Oman.

NDTV Top Stories . 2026-08-08 . ★★★★★

<https://www.ndtv.com/world-news/iran-claims-hormuz-deal-close-but-adds-caveat-for-opening-strait-11884117#publisher=newsstand>

Scale of devastation shocks US city's fire-seasoned residents

QUICK TAKE

Residents of Spokane, a US city long used to fire season, described being shocked by the scale of destruction after the blaze jumped the river into populated neighbourhoods.

BBC World . 2026-08-08 . ★★★★★

https://www.bbc.co.uk/news/articles/cvglpvgljzdo?at_medium=RSS&at_campaign=rss

'I ran because I knew I would die': Russian drones target medics in Ukraine

QUICK TAKE

Ukrainian emergency workers say Russian drones are increasingly targeting medics, hospitals and ambulances, with paramedics describing having to run from strikes on rescue sites.

BBC World . 2026-08-08 . ★★★★★

https://www.bbc.co.uk/news/articles/c3r073eqvrjo?at_medium=RSS&at_campaign=rss

You are caught up. Now stay ahead.

Tomorrow's PDF lands at 00:00 IST - the exact stack you need to stay one day ahead of everyone else in AI.